Comments on Schiffer on Vagueness

Cian Dorr

24 April 2018

1. Some claims Stephen is committed to

(*Background:* Tom utters the sentence 'A boy was here a little while ago'.)

- There is no proposition p such that Tom meant p
- There is no area α such that Tom referred to α by uttering 'here'; there is no period of time π such that Tom referred to π by uttering 'a little while ago'; there is no property ϕ such that Tom referred to (expressed) ϕ by uttering 'boy'.
- There is no area α such that the relevant token of 'here' referred to α; there is no period of time π such that the relevant token of 'a little while ago' referred to π; there is no property ϕ such that the relevant token of 'boy' referred to ϕ. Similarly, there is no area, period of time, or property that is the *semantic value* of any of these tokens.
- There is no area α and period of time π such that Tom meant that a boy was in α during π; there is no property ϕ, area α and period of time π such that Tom meant that an instance of ϕ was in α during π. Similarly: if Tom just said 'I live here', or 'I live there' (pointing on a map), or (extrapolating) 'I live in Midtown', there's no area such that he meant that he lived in it.
- There is no area α such that Tom *intended to refer* to α, or believed he could refer to α, or believed that there was once a boy in α, or believed …α…, or intended …α…, or….

But this is far from being a blanket error theory about speech and attitude reports! For example:

> We surely do want to say that an utterance of a sentence such as 'In uttering "Mon professeur de philosophie est léthargique" Pierre said, and thus meant, that his professor of philosophy was lethargic' might well be true. (13)
>
> it's true that what led to his utterance was his knowing that a boy had been in his vicinity a little while before he spoke (14)

2. Things and precise things

There are some related claims containing the word 'precise' that Stephen regards as *obviously* true. For example, he says that

> it ought to be obvious that there was no precise area to which [Tom] intended to refer when he produced [his utterance], nor would he have thought there was any need to refer to such an area. (6)

A 'precise area' is one that has 'precise conditions of individuation', where I take it this is a proposition; a precise proposition is one such that 'necessarily there is a fact of the matter as to what truth-value, if any, it has'.

Given that ‘precise’ is being used in this technical sense, I don’t think anything stated in terms of it is obvious. (Fortunately Stephen also has *arguments* which we’ll get to later; here I’m just insisting that an argument is needed.)

- Is the claim at least obvious when ‘precise’ is used in the ordinary, non-technical sense? Maybe, but I think we need to be super careful even in accepting claims involving that as obvious, since there are super-quick arguments from some seemingly obvious claims to super-controversial conclusions, such as the following:

  > There is no precise number of hairs such that having exactly that number of hairs does not entail being bald, while having less that that number of hairs does entail being bald. But every number is a precise number. So, there is no number of hairs such that having exactly that number of hairs does not entail being bald, while having less that that number of hairs does entail being bald.

  Note that only classicists like me but prominent non-classicists like Field reject this conclusion. Moral: let’s steer clear of this folk use of ‘precise’ to (seemingly) apply to non-linguistic items.

3. Two motivating thoughts: uniqueness and abundance

So, on to the arguments. In what follows let F abbreviate one of the expressions for which Stephen is defending ‘There is no F’, e.g. ‘area that Tom referred to by saying “here”’. In the cases where the argument for one of these eliminativist claims doesn’t just rest on another eliminativist claim, I think it’s clear that key premises include the following:

Uniqueness: There is at most one F
Abundance: For any x, if x is F, there are enormously many things that are extremely similar to x in all underlying respects.

- What are “underlying respects”? We could imagine trying to sharpen this up by reference canonical language — maybe it should include not only the language of physics but also modal operators, a ‘definitely’ operator, etc. But I’ll leave it schematic.
- Note that the “underlying respects” can include *non-qualitative* respects. For areas, being “similarity in underlying respects” would include being such that the volume of the overlap is much greater than the volume of the symmetric difference.

OK; so what further premise can we add to get from *Uniqueness* and *Abundance* to ‘There is no F’? Here’s a possibility:

Tolerance: There are no x and y such that x is extremely similar to y in underlying respects, x is F, and y is not F.

or less committally:

Weak Tolerance: There are no x and y such that x is extremely similar to y and many other things in underlying respects, x is uniquely F, and y is not F.

The argument from *Uniqueness*, *Abundance*, and *Tolerance* (or *Weak Tolerance*) to 'There is no F' is valid in classical logic, which is good enough for me.[1]

- But does Stephen actually accept this kind of reasoning? It's not so clear to me: there are some passages in the paper that suggest that he must be willing to resist some arguments not a million miles from this one. For example:

  > If there were a real number j that was the largest number such that every token of 'boy' had to be true of every human male whose age in milliseconds ≤ j, then… But of course there can be no such number… as j… (22)

  Taken at face value (and generalized from real to natural numbers) this looks to require rejecting the argument from '*n* is F' and '*n* is not the largest F number' to '*n*+1 is F'.[2] A view that allowed this might be expected to also reject the argument from *Uniqueness*, *Abundance*, and (*Weak*) *Tolerance* to 'There is no F'. On the other hand, such a view would probably also reject *reductio* in general, whereas Stephen often seems to be arguing for not-P based on deriving (allegedly) absurd consequences from P.

4. Background: Unger's argument that there are no people

Back in the 70s, Peter Unger defended arguments like the one you'd get by plugging 'cat in my apartment' or 'person talking to you right now' into *Uniqueness*, *Abundance*, and *Weak Tolerance*. He concluded that there are no cats, no people, etc.

- Note that these arguments are not in any way *about* language, though of course reflection on language might help undermine their appeal, diagnose what's wrong with them, or alternatively bolster their premises.

I don't think Stephen's on board with this. So it's important to see whether Stephen's case for his instantiations of the argument-form can be pulled apart from Unger's case for his instantiations.

In some instantiations, the difference could be in *Abundance*. For example, maybe Nibbles cat does not count as "extremely similar in underlying respects" to anything whose spatial boundary differs from Tibbles's by just one hair on the grounds that Tibbles is conscious and none of those other things are conscious. But there are serious limits to how far this sensibility can take you in resisting Unger-style arguments.

5. How to get from *Uniqueness* and *Abundance* to 'There is no F'?

Stephen appeals to several different kinds of considerations when he is arguing from the relevant instances of *Uniqueness* and *Abundance* to 'There is no F'. Some of the arguments push the case for some instances of F back to the case of other instances of F; For example, he argues

---

[1] Suppose there is an F; call it *a*. By *Abundance* there is something distinct from but extremely similar to *a*; call it *b*. Since by *Uniqueness* every F is identical to *a*, *b* is not F, contradicting *Tolerance*. So by *reductio*, there is no F. This is also intuitionistically valid.

[2] Kit Fine's "compatibilist logic" might be able to block this reasoning be blocked (because conjunctive syllogism

that there's no area that's the semantic value of the relevant token of 'here' partly on the grounds that there's no area that Tom refers to with that token, and argues for this partly on the grounds that there is no area Tom *intended* to refer to with 'here'.[3] So let's focus on the starting points. Here I see several strategies.

The appeal to the need for an explanation

The first strategy is a general demand for explanation: the thought is that for some x rather than y to be (uniquely) F, there would have to be some *explanation* for x's being F and y's not being F, but when x and y are similar in underlying respects there can be no such explanation:

> In other words (more or less), there is nothing to explain how just one of indefinitely many precise areas could have had a feature not possessed by any of the other areas that would have made it alone the area to which Tom referred. (7)

- Problem: however compelling this is, it doesn't seem any more compelling for Stephen's values of F than for Unger's. It's equally hard to think of anything sensible to say in answer to the challenge to explain how just one of the many extremely similar cat-candidates could have a feature not possessed by any of the others that would have made it alone a cat.

The appeal to indistinguishability

In discussing areas, Stephen sometimes appeals to the idea that the extremely similar candidates are *perceptually indistinguishable* to the relevant person (Tom). This is taken to support the conclusion that for each candidate α, Tom 'will have no perceptual way of thinking of α that would do the job' of enabling Tom to think anything about α, including in particular intending to refer to α with 'here' (7). And Stephen takes *perceptual* ways of thinking to be the best candidates, since 'I am unable to conceive of any knowledge by description of α he might have that would enable him to intend to refer to α under it'.

- What does 'indistinguishability' mean here? And is it true that if x is indistinguishable from y to Tom, Tom cannot have a perceptual way of thinking about x?
- One proposal: 'x is indistinguishable from y to Tom' means something like 'Tom is not in a position to know that $x \neq y$', or 'Tom is not in a position to know that x is not exactly similar to y'. The problem is that this might be true just because Tom is not even in a position to *think* that $x \neq y$ (because he can't simultaneously have *de re* beliefs about x and y): it's hard to see why this would pose any obstacle to thinking perceptually about x.

The appeal to knowledgeable uptake

> We refer to things in order to make known to our hearers what we are talking about. In a normal case, such as Tom's, a speaker can't refer to a thing if she knows that her hearer won't be able to know to what she was referring. Tom would know that even if there were a precise location to which he wanted to refer, his hearer would have no way of

[3] I'm not sure whether the premise that referring is an 'intentional act' is true in the ordinary sense of 'refer'—consider 'I unintentionally referred to the war'.

knowing which of the indefinitely many eligible precise areas was the one to which he was referring. Given that, he couldn't have intended to refer to any precise area. (7)

The theory that you can't refer to x if you know your hearer won't be able to know you referred to x might be resisted. But the conclusion that there was nothing that Tom's hearer knew Tom was referring to by using 'here' will already be quite radical if it generalizes to other cases where *Abundance* is in play.

I think the basic thought here is something like this:

> If (*per impossibile*) unique reference to or thought about one of the many candidates was feasible, it would be very *fragile*—given the similarity of the candidates, anyone who managed to single out one of them could very easily have singled out a different one. So it would be quite *unlikely* for two people to have thoughts about the same one, as would have to happen for there to be some α such that the speaker truly believed that Tom intended to refer to α. So even if by luck this did happen, the true belief in question would not count as knowledge.

Never mind the last step: the conclusion that true belief is unlikely in such cases is already pretty radical.

- The argument to that step will be resisted by 'social externalists' like Williamson. According to them, there are strong "constitutive connections" between different people's thought-contents, so even if it's unlikely for anyone to have thoughts about a certain α, conditional on one person doing so it'll be likely that others (e.g. their hearers) to do so as well.
    - But I do think that there's an argument in the vicinity that's harder for social externalists to resist, where we focus on differences across nearby possible worlds rather than differences within the community—see Dorr and Hawthorne, 'Semantic Plasticity and Speech Reports'. My own preferred response involves keeping 'There is at least one F' and giving up *Uniqueness*.

6. From areas and periods of time to properties

In the case of properties, Stephen anticipates a certain way of defending the claim that there *is* a unique property that Tom refers to (and intends to refer to, etc.) using the relevant token of 'boy'. This defence involves saying that boyhood (= the property of being a boy) is (context-insensitively) expressed by 'boy' in English, whereas the many very similar properties are not, and that's why English speakers get to think about and refer to boyhood rather than any of its neighbors.

- One challenge you might push against this view is a generic explanatory one: given the extreme similarities between boyhood and all those other properties, what explains the fact that it rather than any of them gets to be the semantic value of 'boy' in this community's language?

- But Stephen doesn't argue like that: instead, he emphasizes (correctly in my view) the phenomenon he calls "penumbral shift": roughly, that pretty much every vague word is context sensitive in the sense that you can mean different things by it, with the result that you can both definitely say something true and definitely say something false by uttering a sentence involving it without meaning anything relevantly different by any other word in that sentence.

7. The upshot for 'compositional semantics'

Stephen concludes that the languages we speak don't have "compositional truth-theoretic semantics", where

> By a compositional truth- theoretic semantics for a language L I mean a finite theory that generates for each of the infinitely many truth-evaluable sentences of L a theorem that specifies, for each truth-value the theory recognizes, what must be the case in order for a token of the sentence to have that truth-value, and, of course, to have it as a function of semantic values of its constituent expression tokens. (26)

The reason is that some sentence-tokens (e.g. Tom's) have truth values even though their constituent expression tokens (e.g. the tokens of 'here' and 'boy) have no semantic values.

Here are some questions we may want to discuss:

- Is Stephen's definition of 'compositional truth theoretic semantics' too narrow, and if so, can it be broadened without undermining the argument? (Note that many theorists don't seem to use the ideology of tokens having semantic values at all.)
- Can Stephen's conclusion that the relevant tokens lack semantic values be resisted by shifting to some sufficiently different conception of what the semantic values of such tokens are like?
- Is there a reasonable and principled way of accepting Stephen's conclusions (in section I) about speaker-meaning, speaker-reference, and speaker-attitudes while rejecting the corresponding conclusions (in section II) about word- and expression-meaning?