Russell contra Impartiality

Cian Dorr

27 February 2023

1. General ethical setup

The topic: comparative *goodness* of worlds and of prospects *with respect to people's well-being*. Not overall goodness; not what we should do.

- Does the evaluation of these abstract prospects come apart from the evaluation of *actions* that have those prospects? If so, how well do we understand it? If not, isn't 'It's the best available action, but you shouldn't do it' incoherent?
- How much value-relevant stuff is hived off by the restriction to "people's well-being"?

2. The thesis of impartiality

Impartiality: A benefit or harm to one individual makes the same difference to how good things are overall as would the same benefit or harm to any other individual.

A precifisification that avoids difficulties associated with comparing the size of benefits and harms between people:

Numbers*: For any worlds u and v, if u and v are just as good for all but n people $i_1,\ldots, i_n$, then for any other n people $j_1, ..., j_n$ there is some world w which is just as good for each individual except $j_1, ..., j_n$, such that w and v are equally good overall.

Numbers* follows from the n=1 case, aka *Compensation*, given the transitivity of *equally good*.

- *Worry:* what if $i_1,\ldots, i_n$ are doing terribly at u and great at v, whereas $j_1, ..., j_n$ are already living wonderful lives at u?
- We could get around this worry by strengthening the antecedent so that we also require that for each k, u is just as good for i_k as it is for j_k. However this won't be acceptable to diehard opponents of interpersonal welfare comparison.

3. The argument from Ex Ante Pareto and Dominance

Oh No!

	1/6	1/6	1/9	1/9	2/27	2/27	…	$2^{n-1}/3^n$	$2^{n-1}/3^n$

Leave	A1 B1	A1 B1	A2:A3 B2:B3	A2:A3 B2:B3	A4:A7 B4:B7	A4:A7 B4:B7	…	A2^{n}:A2^{n+1}–1 B2^{n}:B2^{n+1}–1	A2^{n}:A2^{n+1}–1 B2^{n}:B2^{n+1}–1
Bury	A2:A3	B2:B3	A4:A7	B4:B7	A8:A15	B8:B15	…	A2^{n+1}:A2^{n+2}–1	B2^{n+1}:B2^{n+1}–1

A2:A3 is, e.g., the outcome in which A2 and A3 suffer and everyone else is fine. (We can think of them as people, or generations of people.)

Impartiality implies that in each state, the outcome of *Leave it* and the outcome of *Bury It* are equally good.

Given this, the following principle implies that the two *prospects* are equally good:

Weak Dominance: If in every state (with nonzero probability) the outcome of *X* is at least as good as the outcome of *Y*, then *X* is at least as good as *Y*.

- This is perhaps slightly less sacrosanct than regular

Dominance: If the set of states in which the outcome of *X* is at least as good as the outcome of *Y* has probability 1, and the set of states in which the outcome of *X* is better than the outcome of *Y* has positive probability, then *X* is better than *Y*.

But we can modify the argument to use this principle instead by slightly sweetening some outcome of *Leave It*, so that Dominance+Impartiality will imply that Leave It is better.

By contrast, the following principle implies that *Bury It* is better:

Ex Ante Pareto: If prospect *X* is at least as good as prospect *Y* for every individual and better than *Y* for at least one individual, then *X* is better than *Y*.

(Actually the argument only needs the weaker version where *X* is better than *Y* for everyone.)

For example: A2 has chance 2/9 of suffering given Leave It versus 1/6 given Bury It; A4 has chance 4/27 versus 1/9; A8 has chance 8/81 versus 2/27… (That's why Jeff chose the sequence 1/6 + 1/9 + 2/27… rather than 1/4 + 1/8 + 1/16…)

4. The argument from Ex Ante Separability and Stochastic Dominance

An Alderaan-only choice:

	1/6	1/6	1/9	1/9	2/27	2/27	…	$2^{n-1}/3^{n}$	$2^{n-1}/3^{n}$
Leave-A	A1	A1	A2:A3	A2:A3	A4:A7	A4:A7	…	A2^{n}:A2^{n+1}–1	A2^{n}:A2^{n+1}–1

Bury-A	A2:A3	A4:A7	A8:A15	…	A2^{n+1}:A2^{n+2}–1

Here, in addition to the argument from Ex Ante Pareto,we can give a different argument that Bury is better, appealing to the uncontroversial claim that outcomes where more people suffer are worse, together with

Stochastic Dominance: If for each outcome o, the probability of the set of states in which the outcome of *X* is at least as good as o is at least as great as the probability of the set of states in which the outcome of *Y* is at least as good as o, and for some outcome o, the probability of the set of states in which the outcome of *X* is at least as good as o is greater than the probability of the set of states in which the outcome of *Y* is at least as good as o, then *X* is better than *Y*.

number of sad people	≥1	≥2	≥4	≥8	…
Probability given *Leave*	1	2/3	4/9	8/27	…
Probability given *Bury*	1/2	1/2	1/3	2/9	…

Now consider a case where the decision is made separately for each planet:

	1/6	1/6	1/9	1/9	2/27	2/27	…	$2^{n-1}/3^n$	$2^{n-1}/3^n$
Leave A, leave B	A1 B1	A1 B1	A2:A3 B2:B3	A2:A3 B2:B3	A4:A7 B4:B7	A4:A7 B4:B7	…	A2^n:A2^{n+1}–1 B2^n:B2^{n+1}–1	A2^n:A2^{n+1}–1 B2^n:B2^{n+1}–1
Leave A, bury B	A1	A1 B2:B3	A2:A3	A2:A3 B4:B7	A4:A7	A4:A7 B8:B15	…	A2^n:A2^{n+1}–1	A2^n:A2^{n+1}–1 B2^{n+1}:B2^{n+2}–1
Bury A, leave B	A2:A3 B1	B1	A4:A7 B2:B3	B2:B3	A8:A15 B4:B7	B4:B7	…	B2^n:B2^{n+1}–1	A2^{n+1}:A2^{n+2}–1 B2^n:B2^{n+1}–1
Bury A, Bury B	A2:A3	B2:B3	A4:A7	B4:B7	A8:A15	B8:B15	…	A2^{n+1}:A2^{n+2}–1	B2^{n+1}:B2^{n+1}–1

Note that Leave A, Leave B and Bury A, Bury B are respectively identical to 'Leave' and 'Bury' from the previous section; so given Impartiality+Dominance, they are equally good (and a slightly sweetened version of the former is better).

Pushing in the other direction we have the following principle:

(Ex Ante) Separability: If we divide the universe up into two separated parts A and B, and prospects *X* and *Y* are such that their outcomes are certain to only differ with respect to what happens in A, and similarly for prospects *X′* and *Y′*, and then *X* and *X′* are certain to only differ with respect to what happens in B, and similarly for *Y* and *Y′*, then *X* is at least as good as *Y* iff *X′* is at least as good as *Y′*.

Earlier we argued that Bury A is better than Leave A. By parity of reasoning, Bury B is better than Leave B. But then by Separability, Bury A, Leave B is better than Leave A, Leave B, and Bury A, Bury B is better than Bury A, Leave B. So by transitivity, Bury A, Bury B is better than Leave A, Leave B.

5. Some doubts about Separability

> *Magrathean Dilemma.* The Vogons have recently destroyed the Earth. Fortunately, the Magratheans have constructed not one but two exact duplicates of the Earth as it was just before its destruction, Earth-A and Earth-B. Currently both planets are in stasis; if one or both is activated, its subsequent history will (thanks to determinism in the laws) proceed in all qualitative respects exactly as Earth's would have if it had not been destroyed. A coin will be tossed: whichever side you pick, Earth-A will be activated if the coin lands on that side and otherwise discarded. Before picking a side, you learn that your friend was offered the same choice but for Earth-B, and picked Heads. So you face the following prospects:

	Heads	Tails
Pick Heads	Both planets activated	No planets activated
Pick Tails	Earth-B activated	Earth-A activated

I'm inclined to think you should pick Tails here (even if the coin is pretty strongly biased towards Heads).

By contrast, if we delete Earth-B and your friend from the story, your choice would look like this:

	Heads	Tails
Pick Heads Alone	Earth-A activated	No planets activated
Pick Tails Alone	No planets activated	Earth-A activated

Here, intuitively, the options are equally good if the coin is fair, and *Pick Heads* is better if it's biased towards Heads. (And we can derive this from Stochastic Dominance and the assumption that the outcome where Earth-A is activated is better.)

But since *Pick Heads* and *Pick Heads Alone* differ only in their effects on the universe outside Earth-A and similarly for *Pick Tails* and *Pick Tails Alone*, while *Pick Heads* and *Pick Tails* differ only in their effects on Earth-A and similarly for *Pick Heads Alone* and *Pick Tails Alone*, Separability entails that *Pick Tails* is better than *Pick Heads* only if *Pick Tails Alone* is better than *Pick Heads Alone.*

6. Similar doubts about Ex Ante Pareto

Can *Magrathean Dilemma* also be used to push against Ex Ante Pareto? It's not clear how to apply Ex Ante Pareto as we stated it to prospects that can differ as regards which people ever concretely exist—that would turn on controversial questions about whether life can be better than never being born.

But we can modify *Magrathean Dilemma* to avoid this. Say that the Magratheans actually have seven planets ready to go, Earth-A through Earth-G. By default, each will be activated in a grim state corresponding to a post-apocalyptic dystopia. If the coin lands on the side you pick, Earth-A will instead be activated in the much nicer state Earth was in before its destruction; and thanks to your friend's choice, if the coin lands Tails, Earth-B will be activated in that same nicer state. Somehow, the same people will eventually be born on each planet no matter what you choose.

7. Parfitean concerns about the metaphysics of personal identity

Complementing such examples, there is a more theoretical reason to be sceptical of principles like *Ex Ante Pareto* that give a distinctive ethical role to questions about the transworld identity of persons.

A theme from Part 3 of Parfit's *Reasons and Persons*: questions about personal identity are in some sense metaphysically shallow, and coming to appreciate this tends to undermine ethical views on which something ethically important turns on these questions.

- Parfit focused on personal identity across *time*, but the problem is if anything even worse with personal identity across *worlds*, where the shallow-seeming questions arise routinely (e.g. in connection with IVF) and not just in science-fictional scenarios involving brain bisection, teletransportation, etc.

My favored way of cashing out 'metaphysically shallow' appeals to *semantic plasticity*. At nearby worlds, the meanings of relevant words (e.g. 'person') vary such that different answers to the questions come out true. For example, maybe some linguistic communities care about the identity of the sperm and egg, while others care about the identity of their nuclei.

If 'person' is plastic like this while 'better' is not relevantly plastic, it is hard for principless such as Ex Ante Pareto connecting betterness with transworld personal identity to come out *robustly* true. (We could think that we just got lucky to hit on the meaning for 'person' on which they come out true; but this position seems epistemically unstable.)

> *Person Street:* suppose that at w_1 there is a street that's infinite in both directions; in each house along the street, a sperm and egg come together and develop in the normal way. The person who develops from S and E, and everyone to their west lives a happy life; everyone to their east lives a sad life. At w_2, things start out

the same but then the *nucleus* of each sperm and egg is transplanted in to the sperm and egg one step to the east before it develops. If the meaning of 'person' is such that what matters is literal sperm-egg identity, these worlds have the same people and they are all equally well off, so (Ex Post) Pareto entails that they are equally good. If the meaning of 'person' is such that what matters is nucleus-identity, one person is worse off in w_2 and the rest are equally well off, so Pareto entails that w_1 is better than w_2.

- One could avoid such arguments by claiming that 'better' is also semantically plastic in a co-ordinate way, such that Ex Ante Pareto comes out true in all close variant languages. But this isn't very appealing—assuming it's more virtuous to care about the good than not to, and we do in fact care (somewhat) about the good, this approach is one where we are just lucky to be as virtuous as we are.

8. Further concerns about the individuation of persons across *epistemic* possibilities

Arguably, the "probabilities" that we care about in evaluating prospects are something like credences or epistemic probabilities. If so, even making sense of principles like Ex Ante Pareto requires making sense of identity of persons across *epistemic* possibilities. This brings up even harder questions related to the multiplicity of "guises" under which one can think about a single person.

9. Jeff's alternative to Impartiality: discounted utilitarianism

One way to give up Impartiality would be to appeal to massive amounts of *incomparability*—e.g., we could say that the only way for one prospect to be at least as good as another is for it to be at least as good for everyone. This is radical, and I'm not sure there's any way to do it with milder amounts of incomparability.

Jeff explores a different way to give up Impartiality while preserving all the other principles (and comparability): discounted utilitarianism. Each person has a 'moral importance score', where the sum of the scores is finite. To get the goodness of a world we multiply each person's well-being by their importance, and then add up.

An incautious reader might think that the proposal is that moral importance depends on one's temporal location. But given that the view aims to preserve Ex Ante Pareto, a person's importance must be a *necessary* feature of them; otherwise, we'll get violations of Pareto by boosting a person's well-being but diminishing their importance.

Question: does how important a person is even supervene on their *qualitative* (non-haecceitistic), *natural* (non-moral) properties?

- Reason to think it can't: arguably, there could be infinite collections of people who were all alike with respect to which qualitative natural properties they *necessarily* instantiate. But given supervenience, they would then all have to all alike with respect to their importance.

- If it doesn't supervene, then we will plausibly be led to the view that whether one world is better than another doesn't supervene on the pattern of qualitative, natural, necessary *relations* between the two worlds. We can have a case where w_1 is better than w_2, and w_3 and w_4 are derived from w_1 and w_2 by the same permutation of individuals, but w_3 isn't better than w_4. This strikes me as extremely strange!

Indeed, given certain auxiliary assumptions, each of the arguments against Impartiality that Jeff considers could be turned directly into an argument against the view that betterness supervenes on the qualitative, natural, necessary relations between worlds. One auxiliary premise we'll need is something along the following lines:

Extreme Anti-Essentialism: for any permutation of possible people π there is an associated permutation π^* of possible worlds, such that for any qualitative natural relation R and possible people $i_1 \ldots i_n$, $R(i_1 \ldots i_n)$ at w iff $R(\pi i_1 \ldots \pi i_n)$ at πw.

Using this, we can (e.g.) redo *Oh No!* such that each for relevant pair of worlds where the same number of people suffer, one of the worlds is derived from the other by applying some cyclic permutation. We'll also need to

10. The anti-Pascalian argument

Pascal's Principle No finite gain is as good as any small increase in the probability ofsome infinite good.

Jeff compellingly brings out what a disaster the truth of Pascal's Principle would be for ethics.

Jeff reports that Pascal's Principle follows from (i) Impartiality; (ii) a certain weakening of Stochastic Dominance, and (iii) the following principle:

Ex Post Infinite Pareto If w_1 is better for infinitely many people than w_2 and at least as good for everyone, then w_1 is better than w_2.

- This is a striking result! But even though Ex Post Infinite Pareto is much weaker than Ex Ante Pareto, it seems to me that giving it up shouldn't seem like a massive cost to those of us who share Parfit's scepticism about ethical principles that implicate questions about the transworld identity of persons.